

Codeflix - Calculating Churn Rates

Learn SQL from Scratch

Matt Kerr

Cohort 2018/12/03

A decorative graphic in the bottom right corner consisting of a light blue rounded square overlapping a white triangle.

Table of Contents

1. GET FAMILIAR WITH THE COMPANY.

- How many months has the company been operating? Which months do you have enough information to calculate a churn rate?
- What segments of users exist?

2. WHAT IS THE OVERALL CHURN TREND SINCE THE COMPANY STARTED?

3. COMPARE THE CHURN RATES BETWEEN USER SEGMENTS.

- Which segment of users should the company focus on expanding?

1. GET FAMILIAR WITH THE COMPANY.

We have been provided with customer subscription data from the company Codeflix.

Our goal is to evaluate their monthly churn rates and search for any trends that might inform their growth strategy.

First, let's just take a look at the `subscriptions` table they've given us to get a general sense of how it's structured.

```
1 SELECT *  
2 FROM subscriptions  
3 LIMIT 15;
```

Query Results			
id	subscription_start	subscription_end	segment
1	2016-12-01	2017-02-01	87
2	2016-12-01	2017-01-24	87
3	2016-12-01	2017-03-07	87
4	2016-12-01	2017-02-12	87
5	2016-12-01	2017-03-09	87
6	2016-12-01	2017-01-19	87
7	2016-12-01	2017-02-03	87
8	2016-12-01	2017-03-02	87
9	2016-12-01	2017-02-17	87
10	2016-12-01	2017-01-01	87
11	2016-12-01	2017-01-17	87
12	2016-12-01	2017-02-07	87
13	2016-12-01	Ø	30
14	2016-12-01	2017-03-07	30
15	2016-12-01	2017-02-22	30

1. GET FAMILIAR WITH THE COMPANY. - CONT'D

How many months has the company been operating?

Which months do you have enough information to calculate a churn rate?

1. GET FAMILIAR WITH THE COMPANY. - CONT'D

How many months has the company been operating?

Which months do you have enough information to calculate a churn rate?

```
1 SELECT MIN(subscription_start) AS 'sub_start_min',
2        MAX(subscription_start) AS 'sub_start_max',
3        MIN(subscription_end) AS 'sub_end_min',
4        MAX(subscription_end) AS 'sub_end_max'
5 FROM subscriptions;
```

Query Results			
sub_start_min	sub_start_max	sub_end_min	sub_end_max
2016-12-01	2017-03-30	2017-01-01	2017-03-31

The earliest subscription_start date is 2016-12-01 and the latest subscription_end date is 2017-03-31. Based on the data we have been provided, we can see that **Codeflix has been operating from at least December 1, 2016 until March 31, 2017.**

1. GET FAMILIAR WITH THE COMPANY. - CONT'D

How many months has the company been operating?

Which months do you have enough information to calculate a churn rate?

```
1 SELECT MIN(subscription_start) AS 'sub_start_min',  
2        MAX(subscription_start) AS 'sub_start_max',  
3        MIN(subscription_end) AS 'sub_end_min',  
4        MAX(subscription_end) AS 'sub_end_max'  
5 FROM subscriptions;
```

Query Results			
sub_start_min	sub_start_max	sub_end_min	sub_end_max
2016-12-01	2017-03-30	2017-01-01	2017-03-31

Monthly churn rate is equal to:

(the number of customers who cancel their subscription in a given month)

÷

(the number of subscribed customers at the start of that month)

1. GET FAMILIAR WITH THE COMPANY. - CONT'D

How many months has the company been operating?

Which months do you have enough information to calculate a churn rate?

```
1 SELECT MIN(subscription_start) AS 'sub_start_min',  
2        MAX(subscription_start) AS 'sub_start_max',  
3        MIN(subscription_end) AS 'sub_end_min',  
4        MAX(subscription_end) AS 'sub_end_max'  
5 FROM subscriptions;
```

Query Results			
sub_start_min	sub_start_max	sub_end_min	sub_end_max
2016-12-01	2017-03-30	2017-01-01	2017-03-31

The earliest subscription_start date is 2016-12-01. This means that the denominator of the monthly churn rate equation is zero for all months before January 2017, and so the first month for which monthly churn can be calculated is January 2017.

The latest subscription_start and subscription_end dates are 2017-03-30 and 2017-03-31, respectively. Given the scope of this project, it seems reasonable to assume that data has only provided until 2017-03-31. And so, the last month for which one can calculate churn is March 2017.

There is enough information to calculate monthly churn rates for January 2017 through March 2017.

1. GET FAMILIAR WITH THE COMPANY. - CONT'D

What segments of users exist?

```
1 SELECT DISTINCT segment
2 FROM subscriptions;
```

Query Results	
segment	
	87
	30

There are two different user segments, designated as 87 and 30, included within the subscriptions table data.

2. WHAT IS THE OVERALL CHURN TREND SINCE THE COMPANY STARTED?

```
1 WITH months AS (  
2   SELECT '2017-01-01' AS first_day,  
3         '2017-01-31' AS last_day  
4   UNION  
5   SELECT '2017-02-01',  
6         '2017-02-28'  
7   UNION  
8   SELECT '2017-03-01',  
9         '2017-03-31'),  
10 cross_join AS (  
11   SELECT *  
12   FROM subscriptions  
13   CROSS JOIN months),  
14 status AS (  
15   SELECT id,  
16         first_day AS month,  
17         CASE  
18           WHEN segment == 87  
19             AND subscription_start < first_day  
20             AND (subscription_end >= first_day  
21                  OR subscription_end IS NULL)  
22             THEN 1  
23           ELSE 0  
24         END AS is_active_87,  
25         CASE  
26           WHEN segment == 30  
27             AND subscription_start < first_day  
28             AND (subscription_end >= first_day  
29                  OR subscription_end IS NULL)  
30             THEN 1  
31           ELSE 0  
32         END AS is_active_30,  
33
```

```
34         CASE  
35           WHEN segment == 87  
36             AND subscription_end >= first_day  
37             AND subscription_end <= last_day  
38             THEN 1  
39           ELSE 0  
40         END AS is_canceled_87,  
41         CASE  
42           WHEN segment == 30  
43             AND subscription_end >= first_day  
44             AND subscription_end <= last_day  
45             THEN 1  
46           ELSE 0  
47         END AS is_canceled_30  
48       FROM cross_join),  
49  
50 status_aggregate AS (  
51   SELECT month,  
52         SUM(is_active_87) AS sum_active_87,  
53         SUM(is_active_30) AS sum_active_30,  
54         SUM(is_canceled_87) AS sum_canceled_87,  
55         SUM(is_canceled_30) AS sum_canceled_30  
56   FROM status  
57   GROUP BY 1)  
58  
59 SELECT month,  
60       (1.0 * sum_canceled_87 / sum_active_87) AS churn_87,  
61       (1.0 * sum_canceled_30 / sum_active_30) AS churn_30  
62 FROM status_aggregate  
63 GROUP BY 1;  
64  
65  
66
```

Query Results		
month	churn_87	churn_30
2017-01-01	0.25089605734767	0.0756013745704467
2017-02-01	0.316916488222698	0.0733590733590734
2017-03-01	0.476894639556377	0.116991643454039

The churn rates for user segments 87 and 30 can be calculated nicely using the code that the capstone project guide helps us to create (shown to the left).

But our churn rate analysis can be made more concise and flexible (for an unknown number of different segments) by altering this code, as suggested by the bonus section of the capstone project guide.

2. WHAT IS THE OVERALL CHURN TREND SINCE THE COMPANY STARTED? - CONT'D

```
1 WITH months AS (  
2   SELECT '2017-01-01' AS first_day,  
3         '2017-01-31' AS last_day  
4   UNION  
5   SELECT '2017-02-01',  
6         '2017-02-28'  
7   UNION  
8   SELECT '2017-03-01',  
9         '2017-03-31'),  
10 cross_join AS (  
11   SELECT *  
12   FROM subscriptions  
13   CROSS JOIN months),  
14  
15 status AS (  
16   SELECT id,  
17         first_day AS month,  
18         segment,  
19         CASE  
20           WHEN subscription_start < first_day  
21             AND (subscription_end >= first_day  
22                OR subscription_end IS NULL)  
23             THEN 1  
24           ELSE 0  
25         END AS is_active,  
26         CASE  
27           WHEN subscription_end >= first_day  
28             AND subscription_end <= last_day  
29             THEN 1  
30           ELSE 0  
31         END AS is_canceled  
32   FROM cross_join)  
33
```

```
34 SELECT segment,  
35        month,  
36        (1.0 * SUM(is_canceled) / SUM(is_active)) AS churn  
37 FROM status  
38 GROUP BY 1, 2;  
39
```

Query Results

segment	month	churn
30	2017-01-01	0.0756013745704467
30	2017-02-01	0.0733590733590734
30	2017-03-01	0.116991643454039
87	2017-01-01	0.25089605734767
87	2017-02-01	0.316916488222698
87	2017-03-01	0.476894639556377

Using this improved code, we can easily calculate the monthly churn rates for any number of different segments. Similar to how we were previously grouping our results by month, we are now grouping our results by both month *and* segment.

2. WHAT IS THE OVERALL CHURN TREND SINCE THE COMPANY STARTED? - CONT'D

```
1 WITH months AS (  
2   SELECT '2017-01-01' AS first_day,  
3         '2017-01-31' AS last_day  
4   UNION  
5   SELECT '2017-02-01',  
6         '2017-02-28'  
7   UNION  
8   SELECT '2017-03-01',  
9         '2017-03-31'),  
10 cross_join AS (  
11   SELECT *  
12   FROM subscriptions  
13   CROSS JOIN months),  
14  
15 status AS (  
16   SELECT id,  
17         first_day AS month,  
18         segment,  
19         CASE  
20           WHEN subscription_start < first_day  
21             AND (subscription_end >= first_day  
22                OR subscription_end IS NULL)  
23             THEN 1  
24           ELSE 0  
25         END AS is_active,  
26         CASE  
27           WHEN subscription_end >= first_day  
28             AND subscription_end <= last_day  
29             THEN 1  
30           ELSE 0  
31         END AS is_canceled  
32   FROM cross_join)  
33
```


```
34 SELECT month,  
35        (1.0 * SUM(is_canceled) / SUM(is_active)) AS churn  
36 FROM status  
37 GROUP BY 1;
```

Query Results

month	churn
2017-01-01	0.16140350877193
2017-02-01	0.188832487309645
2017-03-01	0.27164416203336


By removing the `segment` column from the final `SELECT` statement, and then grouping by only the `month` column, we can also look at Codeflix's monthly churn rates across all user segments. This would be useful for getting an overall sense of how well the subscription service is retaining users.

2. WHAT IS THE OVERALL CHURN TREND SINCE THE COMPANY STARTED? - CONT'D




By plotting the results of the previous two queries, we can see that there is a general upward trend in the monthly churn rates, for all users overall, as well as for each segment (30 and 87) individually.

3. COMPARE THE CHURN RATES BETWEEN USER SEGMENTS.



We can also see that the churn rate for segment 87 is more than 3 times the churn rate for segment 30, for each of the three months evaluated.

3. COMPARE THE CHURN RATES BETWEEN USER SEGMENTS. - CONT'D



Which segment of users should the company focus on expanding?

Codeflix appears better able to retain users of segment 30, compared to users of segment 87, as indicated by segment 30's considerably lower monthly churn rates.

Therefore, Codeflix should focus on expanding the number of users in segment 30. Of course, that is just the simple answer. It's the conclusion, I assume, that the creators of this capstone project had hoped we could reach by the end of this course.

BUT...

3. COMPARE THE CHURN RATES BETWEEN USER SEGMENTS. - CONT'D

Which segment of users should the company focus on expanding? - CONTINUED

The truth is, we don't know anything about what differentiates user segments 30 and 87. Was this an A/B test? What sort of variable was tested? One possibility is that these two segments represent two different customer demographics, in which case the focused expansion of segment 30 seems like reasonable advice, at first.

But in the real world, **additional metrics** (such as user acquisition cost, segment market size, and average revenue per user) and **potential remedies** (like demographic/family discounts) would need to be considered before making such consequential recommendations.

3. COMPARE THE CHURN RATES BETWEEN USER SEGMENTS. - CONT'D

Which segment of users should the company focus on expanding? - CONTINUED

For instance, what if the cost to acquire users for segment 87 is 1/10th the cost to acquire users for segment 30?

And what if segment 87 users are primarily young college students with minimal income? Perhaps by lowering the subscription cost slightly for students (if they sign up with .edu email addresses) for a 4-year term, Codeflix can better retain and garner loyalty from this segment.

And what if the segment 87 market size is 100 times greater than the segment 30 market size?

If it no longer seems as obvious that Codeflix should focus on expanding segment 30, then I have made my point. But to drive it home, here it is in fewer words...

3. COMPARE THE CHURN RATES BETWEEN USER SEGMENTS. - CONT'D

Which segment of users should the company focus on expanding? - CONTINUED

Churn rate is just one piece of the puzzle and should be used in conjunction with other available user and market metrics to develop a robust and successful growth strategy.